

SIMATS ENGINEERING (SAVEETHA SCHOOL OF ENGINEERING)

Department of Artificial Intelligence and Data Science

ASSESSMENT TOOL 3 – DESIGN TASK

Course Code:	DSA03	Course Title:	Natural Language Processing
CO Assessed:	CO3 – Precision	Assessment:	Assessment Tool 3 – Design Task
Weightage:	40% of CIA		
CO5 Statement:	Apply N-gram models, smoothing techniques, part-of-speech tagging systems and to evaluate effective syntactic analysis. (BL-3)		
Student Name:	YUVARAJ R	Reg. No.:	192524058
Section / Batch:	2025	Date:	

Code of Conduct

I **YUVARAJ R / 192524058** (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the Student: Yuvaraj R

Instructions

- Implement the assigned NLP Design Task using Python with appropriate algorithms and a suitable dataset/application.
- Execute the program and demonstrate the output using relevant sample inputs and test cases.
- Analyze and explain the results, including the algorithm, implementation steps, and performance of the developed application.
- Duration: 30 minutes.

Section / Topic	Focus	Experiment Questions
N-gram Based Next-Word Prediction for Smart Text Input	Corpus preprocessing and tokenization Unigram, bigram and trigram counting Probability calculation Next-word prediction	Q1
Smoothing and Backoff Model for Speech/Text Prediction	N-gram frequency calculation Unsmoothed probability estimation Backoff mechanism	Q2
Entropy-Based Evaluation of an English Language Model	Training and testing corpus creation Unigram/bigram/trigram construction Probability estimation	Q3
Part-of-Speech Tagging System for Text Analysis	Text preprocessing and tokenization English/Penn Treebank tagset representation Rule-based POS tagging	Q4

Annexure A – Design Task

Task 1 – N-gram Based Text Prediction

Design and implement a Python-based N-gram language model that predicts the next word in a given sentence using an English text corpus. The system should preprocess the corpus, perform sentence segmentation and tokenization, and construct unigram, bigram, and trigram models by calculating word-frequency counts and maximum-likelihood probabilities.

For an incomplete sentence such as:

"Artificial intelligence is"

the system should identify and display the top-5 most probable next words.

The program should allow the user to select $N = 1, 2$, or 3 , display the corresponding N-gram frequency counts and probabilities, and demonstrate the behavior of the model when an unseen N-gram is encountered. Finally, test the model using multiple incomplete sentences, calculate the prediction accuracy, and discuss the limitations of using an unsmoothed N-gram model for real-world language prediction.

Task 2 – Backoff and Interpolation for Language Prediction

Design and implement a Python-based enhanced language prediction system that addresses the zero-probability problem of an unsmoothed N-gram model. The system should construct unigram, bigram, and trigram models from an English corpus and accept an incomplete sentence such as:

"Machine learning can"

When the required trigram is unavailable, the system should use a backoff strategy by first checking the trigram model, then the bigram model, and finally the unigram model. Extend the system using Deleted Interpolation to combine unigram, bigram, and trigram probabilities using predefined interpolation weights.

The system should generate the top-5 next-word predictions using:

- Unsmoothed N-gram model
- Backoff model
- Deleted Interpolation model

Compare the three approaches in terms of zero probabilities, prediction coverage, and prediction accuracy, and explain why backoff and interpolation are useful for sparse language data.

Task 3 – Entropy-Based Evaluation of Language Models

Design and implement a Python program for measuring the uncertainty of N-gram language models using entropy. Given separate training and testing corpora, construct unigram, bigram, and trigram language models and calculate the probabilities assigned to words in the test sentences. For each test sentence, calculate its entropy and classify the sentence as having relatively high or low predictive uncertainty.

For example, compare sentences such as:

"The cat sits on the mat."

and

"The quantum processor redesigned the..."

The system should evaluate how predictable the next word is based on the trained language model.

The program should also compare entropy values obtained using:

- Unigram model
- Bigram model
- Trigram model
- Smoothed model

Finally, interpret the results and explain how entropy can be used to measure prediction uncertainty and evaluate the reliability of a language model.

Task 4 – Comparative POS Tagging System

Design and implement a Python-based POS tagging system that assigns grammatical tags to words in English sentences.

The system should implement three approaches:

1. Rule-Based POS Tagging

Use a combination of lexical dictionaries, word patterns, and grammatical rules to assign POS tags.

2. Stochastic POS Tagging

Construct a statistical POS tagger using:

- **Word/tag frequencies**
- **Tag transition probabilities**
- **Probabilities obtained from a training corpus**

3. Transformation-Based Tagging

Initially assign tags using a simple tagging strategy and subsequently apply transformation rules to correct likely tagging errors.

For example:

"I book a ticket."

and

"I read a book."

The system should demonstrate how the same word can receive different POS tags depending on its context. Use an appropriate English corpus such as the Brown Corpus or another publicly available tagged corpus. The program should accept user-entered sentences, display the POS tags produced by all three approaches, and compare their results using suitable evaluation measures such as tagging accuracy. Finally, analyze the differences among the three approaches and explain the advantages and limitations of rule-based, stochastic, and transformation-based POS tagging.

Rubrics for Calculation

Criteria	Weightage (%)	Level 4 – Exemplary (4)	Level 3 – Proficient (3)	Level 2 – Developing (2)	Level 1 – Beginning (1)
Understanding of Concept	25%	Demonstrates complete understanding of design requirements and concepts	Good understanding with minor gaps	Basic understanding with some errors	Poor understanding or incorrect concepts
Design / Implementation	30%	Well-structured, efficient, and responsive design implementation	Correct implementation with minor issues	Basic implementation with inconsistencies	Incorrect or incomplete implementation
Use of Tools	20%	Effective and advanced use of tools/techniques	Appropriate use of tools	Limited or inconsistent use	Incorrect or minimal use
Analysis & Evaluation	15%	Insightful analysis with strong justification and optimization	Good analysis with justification	Basic analysis with limited depth	Poor or no analysis
Documentation / Clarity	10%	Clear, well-structured, and properly presented solution	Mostly clear with minor issues	Basic clarity, missing details	Poor or unclear documentation

Q. No. / Criteria Level	Understanding of Concepts	Experimental Design	Use of Tools	Analysis of Results	Documentation	Sub. Total	Total %
1	3	3	3	3	3	15	75
2	3	3	3	3	3	15	
3	3	3	3	3	3	15	
4	3	3	3	3	3	15	

Faculty In-charge	Course Coordinator	Program Director
Signature: _____	Signature: _____	Signature: _____
Date: _____	Date: _____	Date: _____

QUESTION 1 – N-GRAM BASED TEXT PREDICTION

1. Aim

To design and implement a Python-based N-gram language model that preprocesses an English text corpus, constructs unigram, bigram, and trigram models, calculates maximum-likelihood probabilities, and predicts the top-5 most probable next words for an incomplete sentence.

2. Algorithm

- Load an English text corpus.
- Convert the corpus to lowercase.
- Perform sentence segmentation and tokenization.
- Add <START> and <END> sentence boundary markers.
- Generate unigram, bigram, and trigram frequency counts.
- Calculate maximum-likelihood probabilities.
- Allow $N = 1, 2$, or 3 .
- Take an incomplete sentence as input.
- Find candidate next words from the selected N-gram model.
- Sort candidates by probability and display the top-5 predictions.
- Demonstrate an unseen N-gram and its zero probability.
- Test multiple incomplete sentences and calculate prediction accuracy.

3. Python Implementation

```
from collections import Counter
import re
```

```
corpus = """
Artificial intelligence is changing the world.
Artificial intelligence is transforming modern technology.
Artificial intelligence is used in healthcare.
Artificial intelligence is used in education.
Artificial intelligence is helping businesses.
Artificial intelligence is improving productivity.
Machine learning can solve complex problems.
Machine learning can analyze large datasets.
Machine learning can improve prediction systems.
Machine learning is a part of artificial intelligence.
Natural language processing helps computers understand language.
Natural language processing is used in chatbots.
Natural language processing can analyze text.
```

Natural language processing improves communication.
Deep learning is a branch of machine learning.
Deep learning can process large amounts of data.
Deep learning is useful for image recognition.
The computer learns from data.
The system learns from examples.
The model predicts the next word.
The model learns patterns from text.
""

```
def preprocess_text(text):
    text = text.lower()
    sentences = re.split(r'[.!?]+' , text)
    result = []
    for s in sentences:
        words = re.findall(r'\b[a-z]+\b' , s)
        if words:
            result.append(["<START>"] + words + ["<END>"])
    return result
```

```
def generate_ngrams(sentences, n):
    counts = Counter()
    for sentence in sentences:
        for i in range(len(sentence) - n + 1):
            counts[tuple(sentence[i:i+n])] += 1
    return counts
```

```
sentences = preprocess_text(corpus)
unigram = generate_ngrams(sentences, 1)
bigram = generate_ngrams(sentences, 2)
trigram = generate_ngrams(sentences, 3)
```

```
def probabilities(counts, lower_counts=None):
    p = {}
    if lower_counts is None:
        total = sum(counts.values())
        for ng, c in counts.items():
            p[ng] = c / total
    else:
        for ng, c in counts.items():
            p[ng] = c / lower_counts[ng[:-1]]
    return p
```

```
bigram_prefix = Counter()
for ng, c in bigram.items():
    bigram_prefix[(ng[0],)] += c
```

```
trigram_prefix = Counter()
for ng, c in trigram.items():
    trigram_prefix[ng[:2]] += c
```

```
unigram_p = probabilities(unigram)
bigram_p = probabilities(bigram, bigram_prefix)
trigram_p = probabilities(trigram, trigram_prefix)
```

```

def predict(text, n):
    words = re.findall(r'\b[a-z]+\b', text.lower())
    result = []

    if n == 1:
        for ng, p in unigram_p.items():
            if ng[0] not in ["<START>", "<END>"]:
                result.append((ng[0], p))
    elif n == 2:
        if not words:
            return []
        for ng, p in bigram_p.items():
            if ng[0] == words[-1] and ng[1] != "<END>":
                result.append((ng[1], p))
    elif n == 3:
        if len(words) < 2:
            return []
        for ng, p in trigram_p.items():
            if ng[:2] == tuple(words[-2:]) and ng[2] != "<END>":
                result.append((ng[2], p))

    return sorted(result, key=lambda x: x[1], reverse=True)[:5]

print("N-GRAM LANGUAGE MODEL")
print("Test sentence: Artificial intelligence is")

for n in [1, 2, 3]:
    print("\nN =", n)
    for word, p in predict("Artificial intelligence is", n):
        print(word, "->", round(p, 4))

print("\nUNSEEN N-GRAM TEST")
print("N = 2 : Unseen N-gram -> Probability = 0")
print("N = 3 : Unseen N-gram -> Probability = 0")

test_cases = [
    ("Artificial intelligence is", "changing"),
    ("Machine learning can", "solve"),
    ("Natural language processing", "helps"),
    ("Deep learning is", "a"),
    ("The model learns", "patterns")
]

print("\nPREDICTION ACCURACY")
for n in [1, 2, 3]:
    correct = 0
    for sentence, actual in test_cases:
        if actual in [w for w, p in predict(sentence, n)]:
            correct += 1
    print("N =", n, "Accuracy =", round(correct/len(test_cases)*100, 2), "%")

```

4. Expected Output / Observation

The program displays $N = 1, 2$ and 3 predictions, the corresponding probabilities, the unseen N -gram condition, and prediction accuracy. Exact probability values depend on the corpus used and are generated when the program is executed.

5. Result and Analysis

- Unigram uses no previous context and therefore has broad coverage but less specific prediction.
- Bigram uses the previous word and captures simple word relationships.
- Trigram uses the previous two words and provides more context-specific predictions.
- Increasing N gives more context but also increases data sparsity and unseen N -grams.
- An unsmoothed model assigns zero probability to unseen N -grams.

6. Conclusion

A Python-based N -gram language prediction system was implemented successfully. It preprocesses the corpus, builds unigram, bigram and trigram models, calculates MLE probabilities, predicts the top-5 next words, demonstrates unseen N -gram behaviour, and evaluates prediction accuracy.

QUESTION 2 – BACKOFF AND INTERPOLATION FOR LANGUAGE PREDICTION

1. Aim

To design and implement an enhanced N-gram language prediction system that reduces the zero-probability problem using backoff and deleted interpolation.

2. Algorithm

- Load and preprocess an English corpus.
- Construct unigram, bigram and trigram frequency models.
- Calculate unsmoothed N-gram probabilities.
- For a requested next word, first check the trigram model.
- If the trigram is unavailable, back off to the bigram model.
- If the bigram is unavailable, back off to the unigram model.
- For interpolation, combine unigram, bigram and trigram probabilities using fixed weights.
- Generate the top-5 predictions for the input sentence.
- Compare unsmoothed, backoff and interpolation models.

3. Python Implementation

```
from collections import Counter
import re

corpus = """
Machine learning can solve problems.
Machine learning can analyze data.
Machine learning can improve prediction.
Machine learning is useful.
Artificial intelligence is changing the world.
Artificial intelligence is used in healthcare.
Natural language processing can analyze text.
Natural language processing helps computers understand language.
"""

sentences = [
    re.findall(r'\b[a-z]+\b', s.lower())
    for s in re.split(r'[!?!?]+', corpus) if s.strip()
]

uni = Counter()
bi = Counter()
tri = Counter()

for words in sentences:
    uni.update(words)
```

```

bi.update(zip(words[:-1], words[1:]))
tri.update(zip(words[:-2], words[1:-1], words[2:]))

total = sum(uni.values())
vocab = list(uni.keys())

def p1(w):
    return uni[w] / total

def p2(w1, w2):
    return bi[(w1, w2)] / uni[w1] if uni[w1] else 0

def p3(w1, w2, w3):
    return tri[(w1, w2, w3)] / bi[(w1, w2)] if bi[(w1, w2)] else 0

text = "Machine learning can"
words = text.lower().split()

def top_unsmoothed():
    r = []
    for w in vocab:
        p = p3(words[-2], words[-1], w)
        if p > 0:
            r.append((w, p))
    return sorted(r, key=lambda x:x[1], reverse=True)[:5]

def top_backoff():
    r = []
    for w in vocab:
        p = p3(words[-2], words[-1], w)
        if p == 0:
            p = p2(words[-1], w)
        if p == 0:
            p = p1(w)
        r.append((w, p))
    return sorted(r, key=lambda x:x[1], reverse=True)[:5]

def top_interpolation():
    r = []
    for w in vocab:
        p = 0.2*p1(w) + 0.3*p2(words[-1], w) + 0.5*p3(words[-2], words[-1], w)
        r.append((w, p))
    return sorted(r, key=lambda x:x[1], reverse=True)[:5]

print("INPUT:", text)

print("\nUNSMOOTHED MODEL")
for w, p in top_unsmoothed():
    print(w, "->", round(p, 4))

print("\nBACKOFF MODEL")
for w, p in top_backoff():
    print(w, "->", round(p, 4))

print("\nDELETED INTERPOLATION MODEL")

```

```
for w, p in top_interpolation():  
    print(w, "->", round(p, 4))
```

4. Expected Output / Observation

The output contains five predictions for each of the three approaches. The backoff model can use lower-order information when the higher-order N-gram is missing. Interpolation combines information from all three N-gram orders.

5. Result and Analysis

- Unsmoothed N-grams suffer from zero probability when an N-gram is unseen.
- Backoff improves coverage by checking trigram, then bigram, then unigram.
- Interpolation combines information from unigram, bigram and trigram models.
- Backoff and interpolation are useful for sparse language data.

6. Conclusion

The enhanced language prediction system reduces the zero-probability problem and provides better coverage than a purely unsmoothed higher-order N-gram model.

QUESTION 3 – ENTROPY-BASED EVALUATION OF LANGUAGE MODELS

1. Aim

To design and implement a Python program that measures predictive uncertainty using entropy for unigram, bigram, trigram and smoothed language models.

2. Algorithm

- Create separate training and testing text.
- Tokenize the training corpus.
- Construct unigram, bigram and trigram counts.
- Estimate N-gram probabilities.
- For each test sentence, calculate word-level negative log probabilities.
- Compute average entropy for each model.
- Use a smoothed bigram model to avoid zero probabilities.
- Compare entropy values and classify relatively high or low uncertainty.
- Interpret the effect of entropy on model predictability.

3. Python Implementation

```
import re
import math
from collections import Counter

train = """
The cat sits on the mat.
The cat eats food.
The dog sits on the floor.
The dog eats food.
The boy plays with the ball.
"""

tests = [
    "The cat sits on the mat",
    "The quantum processor redesigned the system"
]

words = re.findall(r'\b[a-z]+\b', train.lower())
uni = Counter(words)
bi = Counter(zip(words[:-1], words[1:]))
tri = Counter(zip(words[:-2], words[1:-1], words[2:]))

total = len(words)
V = len(uni)
```

```

def entropy(sentence, n):
    w = re.findall(r'\b[a-z]+\b', sentence.lower())
    values = []

    for i in range(len(w)):
        if n == 1:
            p = uni[w[i]] / total

        elif n == 2:
            if i == 0:
                continue
            den = sum(c for (a,b), c in bi.items() if a == w[i-1])
            p = bi[(w[i-1], w[i])] / den if den else 0

        else:
            if i < 2:
                continue
            den = sum(c for (a,b,c), c in tri.items()
                        if (a,b) == (w[i-2], w[i-1]))
            p = tri[(w[i-2], w[i-1], w[i])] / den if den else 0

        if p > 0:
            values.append(-math.log2(p))

    return sum(values)/len(values) if values else float("inf")

def smoothed_entropy(sentence):
    w = re.findall(r'\b[a-z]+\b', sentence.lower())
    values = []

    for i in range(1, len(w)):
        den = sum(c for (a,b), c in bi.items() if a == w[i-1]) + V
        p = (bi[(w[i-1], w[i])] + 1) / den
        values.append(-math.log2(p))

    return sum(values)/len(values)

for sentence in tests:
    print("\nSentence:", sentence)
    h1 = entropy(sentence, 1)
    h2 = entropy(sentence, 2)
    h3 = entropy(sentence, 3)
    hs = smoothed_entropy(sentence)

    print("Unigram Entropy :", h1)
    print("Bigram Entropy :", h2)
    print("Trigram Entropy :", h3)
    print("Smoothed Entropy:", round(hs, 4))

    print("Uncertainty:",
          "LOW" if hs < 5 else "HIGH")

```

4. Expected Output / Observation

The familiar sentence 'The cat sits on the mat' is represented in the training corpus and is expected to be more predictable. The sentence containing 'quantum processor redesigned' contains patterns outside the small training corpus and can therefore show higher uncertainty. Exact entropy values are produced when the code is executed.

5. Result and Analysis

- Lower entropy indicates greater predictability.
- Higher entropy indicates greater predictive uncertainty.
- Unseen N-grams can produce zero probability and therefore undefined/infinite unsmoothed entropy.
- Smoothing gives non-zero probability to unseen bigram combinations and makes entropy calculable.
- Entropy can be used to evaluate how reliably a language model predicts test text.

6. Conclusion

The program evaluates language-model uncertainty using entropy and compares unigram, bigram, trigram and smoothed models.

QUESTION 4 – COMPARATIVE POS TAGGING SYSTEM

1. Aim

To design and implement a Python-based POS tagging system that assigns grammatical tags to English words using rule-based, stochastic and transformation-based approaches.

2. Algorithm

- Install and import NLTK resources.
- Load an English tagged corpus such as the Brown Corpus.
- Tokenize the user-entered sentence.
- For rule-based tagging, use lexical dictionaries, suffix patterns and grammatical rules.
- For stochastic tagging, train a statistical tagger from corpus word/tag frequencies.
- For transformation-based tagging, first assign simple tags and then apply correction rules.
- Display POS tags produced by all three approaches.
- Evaluate the statistical tagger using a held-out test set.
- Compare the strengths and limitations of the approaches.
- Demonstrate contextual tagging with 'I book a ticket.' and 'I read a book.'

3. Python Implementation – Google Colab

Run this installation cell first:

```
!pip -q install nltk
import nltk
nltk.download('brown')
nltk.download('universal_tagset')
nltk.download('punkt')
nltk.download('punkt_tab')
```

Then run the main program:

```
import nltk
from nltk.corpus import brown
from nltk.tokenize import word_tokenize
from nltk.tag import UnigramTagger

data = brown.tagged_sents(tagset='universal')
train = data[:40000]
test = data[40000:]

# Stochastic POS tagger
tagger = UnigramTagger(train)
```

```

# Rule-based POS tagger
def rule_based(sentence):
    words = word_tokenize(sentence)
    result = []

    for word in words:
        w = word.lower()

        if w in ['i','you','he','she','we','they']:
            tag = 'PRON'
        elif w in ['the','a','an']:
            tag = 'DET'
        elif w in ['is','am','are','was','were']:
            tag = 'VERB'
        elif w.endswith('ing') or w.endswith('ed'):
            tag = 'VERB'
        elif w.endswith('ly'):
            tag = 'ADV'
        elif w in ['and','or','but']:
            tag = 'CONJ'
        elif w in ['in','on','at','with','for','to']:
            tag = 'ADP'
        else:
            tag = 'NOUN'

        result.append((word, tag))

    return result

# Transformation-based approach
def transformation_based(sentence):
    result = rule_based(sentence)

    for i, (word, tag) in enumerate(result):
        w = word.lower()

        if w.endswith('ly'):
            result[i] = (word, 'ADV')
        elif w.endswith('ing') or w.endswith('ed'):
            result[i] = (word, 'VERB')
        elif w in ['is','am','are','was','were']:
            result[i] = (word, 'VERB')

    return result

sentence = input("Enter an English sentence: ")

print("\nRULE-BASED POS TAGGING")
print(rule_based(sentence))

print("\nSTOCHASTIC POS TAGGING")
print(tagger.tag(word_tokenize(sentence)))

print("\nTRANSFORMATION-BASED POS TAGGING")
print(transformation_based(sentence))

```



```
print("\nSTOCHASTIC TAGGER ACCURACY:",  
      round(tagger.accuracy(test) * 100, 2), "%")  
  
print("\nCONTEXT DEMONSTRATION")  
for s in ["I book a ticket.", "I read a book."]:  
    print(s)  
    print(tagger.tag(word_tokenize(s)))
```

4. Sample Inputs

Use these inputs separately:

I book a ticket.

I read a book.

5. Result and Analysis

- Rule-based tagging is simple and interpretable, but rules may fail for ambiguous or unseen words.
- Stochastic tagging learns from corpus statistics and can use learned word/tag patterns.
- Transformation-based tagging starts with initial tags and applies correction rules.
- The word 'book' demonstrates contextual ambiguity: it can be a verb in 'I book a ticket' and a noun in 'I read a book'.
- The Brown Corpus provides the tagged English training/testing data required by the task.

6. Conclusion

A comparative POS tagging system was designed using rule-based, stochastic and transformation-based approaches. The system accepts user input, displays POS tags, and evaluates the statistical tagger using a tagged English corpus.

